

```

from .environment_base import SimpleTrader
import numpy as np
import random
import yfinance as yf

# Ryan's implementation of the environment

class Trader(SimpleTrader):

    def __init__(self, ticker_list, observation_metrics=3, initial_funds=2000, starting_date="2023-04-05", ending_date="2023-10-05"):
        super().__init__(ticker_list, initial_funds=initial_funds, observation_metrics=observation_metrics,
                          starting_date=starting_date, ending_date=ending_date)

        self.epsilon = 1
        self.vix = yf.download(
            "^VIX", start=self.starting_date, end=self.ending_date)
        self.max_portfolio = self.initial_funds

    def reset(self, render=False, seed=None):
        """
        Resets the reinforcement learning environment

        Parameters:
            render (boolean): Whether to display the episodic performance of the
            agent
            seed (int): Seed for predictable behaviour (NOT USED)

        Returns:
            observation (float): The observations for the initial trading day
            info (dict): A dictionary containing additional information about
            the environment (NOT USED)
        """
        if seed != None:
            np.random.seed(seed)

        if render == True:
            self._render_on_completion()

        if not hasattr(self, "funds_history"):
            self.funds_history, self.portfolio_history = [], []
            self.episode_funds, self.episode_portfolio = [], []
            self.render_episodes = False

        self.curr_step = 0
        self.curr_funds = self.initial_funds
        self.portfolio_value = self.initial_funds

        self.num_buys, self.num_sells = 0, 0
        self.buy_percents, self.sell_percents = 0.0, 0.0
        self.owned_shares = np.zeros(self.num_stocks)

        self.returns = []

        self.epsilon = 1
        self.max_portfolio = self.initial_funds

        observation = self._get_observation_ryan_0()
        info = {}

        return observation, info

    def step(self, action_list):
        """
        Driving logic for updating the reinforcement learning environment on
        each trading day

        Parameters:
            action_list (list): A list of actions for each stock in the portfolio

        Returns:
            observation (float): The observations for the current trading day
            reward (float): The reward signal for a given series of actions
            done (boolean): A flag to check whether the current episode has
            terminated
            terminated (boolean): A flag to check whether early stopping of the
            episode has occurred (NOT USED)
            info (dict): A dictionary containing additional information about
            the environment (NOT USED)
        """
        self.curr_step += 1

        done = self.curr_step >= self.num_trading_days

        reward = self._perform_action_ryan_0(action_list)

```

```

if self.render_episodes == True:
    self.episode_funds.append(self.curr_funds - self.initial_funds)
    self.episode_portfolio.append(self.portfolio_value)

    if done:
        self.funds_history.append(self.episode_funds)
        self.portfolio_history.append(
            self.episode_portfolio)
        self.episode_funds, self.episode_portfolio = [], []
        self._get_total_action_count()
        self._get_total_buy_sell_percents()

if not done:
    observation = self._get_observation_ryan_0()
else:
    observation = None

terminated = False
info = {}

return observation, reward, done, terminated, info

def _perform_action_ryan_0(self, action_list):
    """
    Base action function:
    - Retrieves the opening price for a given trading day
    - Performs the actual buy / sell actions from a given action list
    - Updates the current funds, shares held and new portfolio valuation
    - Calculates and returns reward

    Parameters:
        action_list (list): A list of actions for each stock in the portfolio

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    curr_date = self.trading_days[self.curr_step - 1]

    stockVal = 0

    opening_price = [self.stock_data.loc[(
        ticker, curr_date), "Open"] for ticker in self.ticker_list]

    for ii in range(len(self.owned_shares)):
        stockVal += (self.owned_shares[ii] * opening_price[ii])

    buy_reward = 0
    sell_reward = 0

    money_change = []

    for ii, action in enumerate(action_list):

        if action < 0: # sell signal
            max_shares = self.owned_shares[ii]
            num_shares = int(abs(action) * max_shares)

            if self.owned_shares[ii] >= num_shares:
                self.owned_shares[ii] -= num_shares
                money_change.append(num_shares * opening_price[ii])
                self.curr_funds += num_shares * opening_price[ii]

                self.num_sells += 1
                self.sell_percents += abs(action)

            if num_shares == 0:
                # Reduce reward if trying to sell stocks that don't exist
                sell_reward -= abs(action * 10) ** 2

    for ii, action in enumerate(action_list):
        if action > 0: # buy signal
            max_investment = self.curr_funds
            investment = action * max_investment
            num_shares = int(investment / opening_price[ii])
            investment = num_shares * opening_price[ii]

            if self.curr_funds >= investment:
                self.owned_shares[ii] += num_shares
                self.curr_funds -= investment
                money_change.append(investment)

                self.num_buys += 1
                self.buy_percents += action

            if investment == 0:

```

```

        buy_reward -= abs(action * 10) ** 2

# Prevents crashes in the rare event of an action for a stock being equal to 0
for ii, action in enumerate(action_list):
    if action == 0.0: # Hold stock
        money_change.append(0.0)

closing_price = np.array([self.stock_data.loc[(
    ticker, curr_date), "Adj Close"] for ticker in self.ticker_list])

self.portfolio_value = self.curr_funds + \
    sum(self.owned_shares * closing_price)

reward = self._get_reward_ryan_7(
    action_list)

self.previous_portfolio = self.portfolio_value

return reward

def _perform_action_ryan_1(self, action_list):
    """
    Second version of action function:
    - Retrieves the opening price for a given trading day
    - Performs the actual buy / sell actions from a given action list
    - Updates the current funds, shares held and new portfolio valuation
    - Calculates and returns reward
    - Includes epsilon greedy exploration (if below threshold, random actions are taken)

    Parameters:
        action_list (list): A list of actions for each stock in the portfolio

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    epsilon = 0.1

    curr_date = self.trading_days[self.curr_step - 1]

    opening_price = [self.stock_data.loc[(
        ticker, curr_date), "Open"] for ticker in self.ticker_list]

    stockVal = 0

    for ii in range(len(self.owned_shares)):
        stockVal += (self.owned_shares[ii] * opening_price[ii])

    startPortfolio = self.curr_funds + stockVal

    print(f"action: {action_list}")

    buy_reward = 0
    sell_reward = 0

    money_change = []

    x = random.uniform(0, 1)

    if x < epsilon:
        action_list = np.random.uniform(-1, 1, size=(len(action_list),))
        print(f"Random action: {action_list}")

    for ii, action in enumerate(action_list):

        if action < 0: # sell signal
            max_shares = self.owned_shares[ii]
            num_shares = int(abs(action) * max_shares)

            if self.owned_shares[ii] >= num_shares:
                self.owned_shares[ii] -= num_shares
                money_change.append(num_shares * opening_price[ii])
                self.curr_funds += num_shares * opening_price[ii]

                self.num_sells += 1
                self.sell_percents += abs(action)

            if num_shares == 0:
                # Reduce reward if trying to sell stocks that don't exist
                sell_reward -= abs(action * 10) ** 2

        for ii, action in enumerate(action_list):
            if action > 0: # buy signal
                max_investment = self.curr_funds
                investment = action * max_investment
                num_shares = int(investment / opening_price[ii])
                investment = num_shares * opening_price[ii]

```

```

        if self.curr_funds >= investment:
            self.owned_shares[ii] += num_shares
            self.curr_funds -= investment
            money_change.append(investment)

            self.num_buys += 1
            self.buy_percents += action

        if investment == 0:
            # Reduce reward if trying to buy stock without money
            buy_reward -= abs(action * 10) ** 2

# Prevents crashes in the rare event of an action for a stock being equal to 0
for ii, action in enumerate(action_list):
    if action == 0.0: # Hold stock
        money_change.append(0.0)

closing_price = np.array([self.stock_data.loc[(
    ticker, curr_date), "Adj Close"] for ticker in self.ticker_list])

self.portfolio_value = self.curr_funds + \
    sum(self.owned_shares * closing_price)

reward = self._get_reward_ryan_0(
    action_list, buy_reward, sell_reward, money_change, startPortfolio)

self.previous_portfolio = self.portfolio_value

return reward

def _perform_action_ryan_2(self, action_list):
    """
    Third version of action function:
    - Retrieves the opening price for a given trading day
    - Performs the actual buy / sell actions from a given action list
    - Updates the current funds, shares held and new portfolio valuation
    - Calculates and returns reward
    - Includes decaying epsilon greedy exploration to be passed to reward function

    Parameters:
        action_list (list): A list of actions for each stock in the portfolio

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    self.epsilon *= 0.5

    curr_date = self.trading_days[self.curr_step - 1]

    opening_price = [self.stock_data.loc[(
        ticker, curr_date), "Open"] for ticker in self.ticker_list]

    stockVal = 0

    for ii in range(len(self.owned_shares)):
        stockVal += (self.owned_shares[ii] * opening_price[ii])

    startPortfolio = self.curr_funds + stockVal

    print(f"action: {action_list}")

    buy_reward = 0
    sell_reward = 0

    money_change = []

    for ii, action in enumerate(action_list):

        if action < 0: # sell signal
            max_shares = self.owned_shares[ii]
            num_shares = int(abs(action) * max_shares)

            if self.owned_shares[ii] >= num_shares:
                self.owned_shares[ii] -= num_shares
                money_change.append(num_shares * opening_price[ii])
                self.curr_funds += num_shares * opening_price[ii]

                self.num_sells += 1
                self.sell_percents += abs(action)

            if num_shares == 0:
                # Reduce reward if trying to sell stocks that don't exist
                sell_reward -= abs(action * 10) ** 2

    for ii, action in enumerate(action_list):

```

```

        if action > 0: # buy signal
            max_investment = self.curr_funds
            investment = action * max_investment
            num_shares = int(investment / opening_price[ii])
            investment = num_shares * opening_price[ii]

            if self.curr_funds >= investment:
                self.owned_shares[ii] += num_shares
                self.curr_funds -= investment
                money_change.append(investment)

                self.num_buys += 1
                self.buy_percents += action
            if investment == 0:
                # Reduce reward if trying to buy stock without money
                buy_reward -= abs(action * 10) ** 2

# Prevents crashes in the rare event of an action for a stock being equal to 0
for ii, action in enumerate(action_list):
    if action == 0.0: # Hold stock
        money_change.append(0.0)

closing_price = np.array([self.stock_data.loc[(
    ticker, curr_date), "Adj Close"] for ticker in self.ticker_list])

self.portfolio_value = self.curr_funds + \
    sum(self.owned_shares * closing_price)

reward = self._get_reward_ryan_4(
    action_list, buy_reward, sell_reward, money_change, startPortfolio)

self.previous_portfolio = self.portfolio_value

return reward

def _get_reward_ryan_0(self, action_list, buy_reward, sell_reward, money_change):
    """
    Reward function for stock trading environment:
        - Increase portfolio value from initial value

    Parameters:
        action_list (list): A list of actions for each stock in the
        portfolio
        buy_reward (float): Reward for buying stocks
        sell_reward (float): Reward for selling stocks
        money_change (list): A list of the amount of money gained / lost

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    if self.curr_step != len(self.trading_days):
        next_date = self.trading_days[self.curr_step]
    else:
        next_date = self.trading_days[self.curr_step - 1]

    next_opening_price = [self.stock_data.loc[(
        ticker, next_date), "Open"] for ticker in self.ticker_list]

    stockVal = 0
    for ii in range(len(self.owned_shares)):
        stockVal += (self.owned_shares[ii] * next_opening_price[ii])
    next_portfolio = self.curr_funds + stockVal

    reward = (100 * (next_portfolio - self.portfolio_value) / next_portfolio)

    curr_date = self.trading_days[self.curr_step - 1]

    opening_price = [self.stock_data.loc[(
        ticker, curr_date), "Open"] for ticker in self.ticker_list]

    for ii, action in enumerate(action_list):
        reward += (money_change[ii] * action * (next_opening_price[ii] -
            opening_price[ii]) / next_opening_price[ii]) / 10

    reward += (buy_reward + sell_reward)

    return reward

def _get_reward_ryan_1(self, action_list, buy_reward, sell_reward, money_change):
    """
    Reward function for stock trading environment:
        - Increase portfolio value from initial value
        - Decaying Epsilon greedy exploration (gradually shrinking reward
        component to encourage a learned strategy)

    Parameters:

```

```

        action_list (list): A list of actions for each stock in the
        portfolio
        buy_reward (float): Reward for buying stocks
        sell_reward (float): Reward for selling stocks
        money_change (list): A list of the amount of money gained / lost

Returns:
    reward (float): The reward signal for a given series of actions
    """

    if self.curr_step != len(self.trading_days):
        next_date = self.trading_days[self.curr_step]
    else:
        next_date = self.trading_days[self.curr_step - 1]

    next_opening_price = [self.stock_data.loc[(
        ticker, next_date), "Open"] for ticker in self.ticker_list]

    stockVal = 0
    for ii in range(len(self.owned_shares)):
        stockVal += (self.owned_shares[ii] * next_opening_price[ii])
    next_portfolio = self.curr_funds + stockVal

    reward = (100 * (next_portfolio - self.portfolio_value) / next_portfolio)

    curr_date = self.trading_days[self.curr_step - 1]

    opening_price = [self.stock_data.loc[(
        ticker, curr_date), "Open"] for ticker in self.ticker_list]

    for ii, action in enumerate(action_list):
        reward += (money_change[ii] * action * (next_opening_price[ii] -
            opening_price[ii]) / next_opening_price[ii]) / 10

    reward += (buy_reward + sell_reward)

    reward += (self.epsilon * 10)

    return reward

def _get_reward_ryan_2(self):
    """
    Reward function for stock trading environment:
        - Calculate daily return (ratio of portfolio value to previous portfolio
        value)
        - Calculate short-term Sharpe ratio over 20 day window

    Parameters:

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    SHARPE_WINDOW = 20

    daily_return = (self.portfolio_value -
        self.previous_portfolio) / self.previous_portfolio

    self.returns.append(daily_return)

    if len(self.returns) < SHARPE_WINDOW:
        reward = 0
    else:
        ret = np.array(self.returns[-SHARPE_WINDOW:])
        sharpe_ratio = np.sqrt(SHARPE_WINDOW) * np.mean(ret) / np.std(ret)

        reward = sharpe_ratio

    return reward

def _get_reward_ryan_3(self):
    """
    Reward function for stock trading environment:
        - Calculate daily return (ratio of portfolio value to previous portfolio
        value)
        - Calculate short-term Sharpe ratio over 20 day window
        - Dynamic volatility weighting based on VIX index

    Parameters:

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    SHARPE_WINDOW = 20

    daily_return = (self.portfolio_value -

```

```

        self.previous_portfolio) / self.previous_portfolio

self.returns.append(daily_return)

reward = 0

if len(self.returns) >= SHARPE_WINDOW:
    ret = np.array(self.returns[-SHARPE_WINDOW:])
    sharpe_ratio = np.sqrt(SHARPE_WINDOW) * np.mean(ret) / np.std(ret)

    reward += sharpe_ratio

if len(self.returns) > 0:
    curr_volatility = np.std(
        np.array(self.returns[-min(SHARPE_WINDOW, len(self.returns)):]))

    curr_vix = self._get_vix()
    volatility_weight = curr_vix * 0.01

    volatility_penalty = curr_volatility * volatility_weight

    reward -= volatility_penalty

return reward

def _get_reward_ryan_4(self):
    """
    Reward function for stock trading environment:
    - Calculate daily return (ratio of portfolio value to previous portfolio
    value)
    - Calculate short-term Sharpe ratio over 20 day window
    - Dynamic volatility weighting based on VIX index
    - Dynamically adjusted risk aversion (drawdown penalty)

    Parameters:

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    SHARPE_WINDOW = 20

    daily_return = (self.portfolio_value -
                    self.previous_portfolio) / self.previous_portfolio

    self.returns.append(daily_return)
    self.max_portfolio = max(self.max_portfolio, self.portfolio_value)

    reward = 0

    curr_vix = self._get_vix()

    drawdown_threshold = curr_vix * 0.005

    drawdown = (self.max_portfolio - self.portfolio_value) / \
        self.max_portfolio

    if drawdown > drawdown_threshold:
        reward -= drawdown

    if len(self.returns) >= SHARPE_WINDOW:
        ret = np.array(self.returns[-SHARPE_WINDOW:])
        sharpe_ratio = np.sqrt(SHARPE_WINDOW) * np.mean(ret) / np.std(ret)

        reward += sharpe_ratio

    if len(self.returns) > 0:
        curr_volatility = np.std(
            np.array(self.returns[-min(SHARPE_WINDOW, len(self.returns)):]))

        volatility_weight = curr_vix * 0.01

        volatility_penalty = curr_volatility * volatility_weight

        reward -= volatility_penalty

    return reward

def _get_reward_ryan_5(self, action_list):
    """
    Reward function for stock trading environment:
    - Calculate daily return (ratio of portfolio value to previous portfolio
    value)
    - Calculate short-term Sharpe ratio over 20 day window
    - Dynamic volatility weighting based on VIX index
    - Dynamically adjusted risk aversion (drawdown penalty)
    - Trend following based on adjustable weighting

```

```

Parameters:
    action_list (list): A list of actions for each stock in the portfolio

Returns:
    reward (float): The reward signal for a given series of actions
"""

SHARPE_WINDOW = 20
TREND_REWARD = 0.1

daily_return = (self.portfolio_value -
                self.previous_portfolio) / self.previous_portfolio

self.returns.append(daily_return)
self.max_portfolio = max(self.max_portfolio, self.portfolio_value)

reward = 0

curr_vix = self._get_vix()

drawdown_threshold = curr_vix * 0.005

drawdown = (self.max_portfolio - self.portfolio_value) / \
            self.max_portfolio

if drawdown > drawdown_threshold:
    reward -= drawdown

if len(self.returns) >= SHARPE_WINDOW:
    ret = np.array(self.returns[-SHARPE_WINDOW:])
    sharpe_ratio = np.sqrt(SHARPE_WINDOW) * np.mean(ret) / np.std(ret)

    reward += sharpe_ratio

if len(self.returns) > 0:
    curr_volatility = np.std(
        np.array(self.returns[-min(SHARPE_WINDOW, len(self.returns)):]))

    volatility_weight = curr_vix * 0.01

    volatility_penalty = curr_volatility * volatility_weight

    reward -= volatility_penalty

if self.curr_step > 0:
    curr_date = self.trading_days[self.curr_step - 1]

    for ii, action in enumerate(action_list):

        if self.curr_step >= SHARPE_WINDOW:
            prev_date = self.trading_days[self.curr_step - SHARPE_WINDOW]
            stock_price = self.stock_data.xs(
                self.ticker_list[ii], level="Ticker", axis=0).loc[prev_date:curr_date, "Adj Close"].values

            if len(stock_price) > 0:
                ema = np.mean(stock_price)

                curr_price = self.stock_data.loc[(
                    self.ticker_list[ii], curr_date), "Adj Close"]

                if (action > 0 and curr_price > ema) or (action < 0 and curr_price < ema):
                    reward += TREND_REWARD * abs(action)

return

def _get_reward_ryan_6(self, action_list):
    """
    Reward function for stock trading environment:
    - Calculate daily return (ratio of portfolio value to previous portfolio
      value)
    - Calculate short-term Sharpe ratio over 20 day window
    - Dynamic volatility weighting based on VIX index
    - Dynamically adjusted risk aversion (drawdown penalty)
    - Dynamic trend following with adjustable base weight

    Parameters:
        action_list (list): A list of actions for each stock in the portfolio

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    SHARPE_WINDOW = 20
    BASE_TREND_REWARD = 0.1
    TREND_WINDOW = SHARPE_WINDOW

```



```

daily_return = (self.portfolio_value -
                self.previous_portfolio) / self.previous_portfolio

self.returns.append(daily_return)
self.max_portfolio = max(self.max_portfolio, self.portfolio_value)

reward = 0

curr_vix = self._get_vix()

drawdown_threshold = curr_vix * 0.005

drawdown = (self.max_portfolio - self.portfolio_value) / \
            self.max_portfolio

if drawdown > drawdown_threshold:
    reward -= drawdown

if len(self.returns) >= SHARPE_WINDOW:
    ret = np.array(self.returns[-SHARPE_WINDOW:])
    sharpe_ratio = np.sqrt(SHARPE_WINDOW) * np.mean(ret) / np.std(ret)

    reward += sharpe_ratio

if len(self.returns) > 0:
    curr_volatility = np.std(
        np.array(self.returns[-min(SHARPE_WINDOW, len(self.returns)):]))

    volatility_weight = curr_vix * 0.01

    volatility_penalty = curr_volatility * volatility_weight

    reward -= volatility_penalty

if self.curr_step > 0:
    curr_date = self.trading_days[self.curr_step - 1]

    for ii, action in enumerate(action_list):

        if self.curr_step >= SHARPE_WINDOW:
            prev_date = self.trading_days[self.curr_step - SHARPE_WINDOW]
            stock_price = self.stock_data.xs(
                self.ticker_list[ii], level="Ticker", axis=0).loc[prev_date:curr_date, "Adj Close"].values

            if len(stock_price) > 0:
                ema = np.mean(stock_price)

                curr_price = self.stock_data.loc[(
                    self.ticker_list[ii], curr_date), "Adj Close"]

                if len(stock_price) >= TREND_WINDOW:
                    returns = np.diff(stock_price) / stock_price[:-1]
                    volatility = np.std(returns[-TREND_WINDOW:])
                    trend_reward = BASE_TREND_REWARD * (1 + volatility)
                else:
                    trend_reward = BASE_TREND_REWARD

                if (action > 0 and curr_price > ema) or (action < 0 and curr_price < ema):
                    reward += trend_reward * abs(action)

    return reward

def _get_reward_ryan_7(self, action_list):
    """
    Reward function for stock trading environment:
    - Calculate daily return (ratio of portfolio value to previous portfolio
      value)
    - Calculate short-term Sharpe ratio over three different windows with
      different adjustable weighting (short, medium and long)
    - Dynamic volatility weighting based on VIX index
    - Dynamically adjusted risk aversion (drawdown penalty)
    - Dynamic trend following with adjustable base weight

    Parameters:
        action_list (list): A list of actions for each stock in the portfolio

    Returns:
        reward (float): The reward signal for a given series of actions
    """

    SHARPE_WINDOW_SHORT = 20
    SHARPE_WINDOW_MED = 60
    SHARPE_WINDOW_LONG = 100
    BASE_TREND_REWARD = 0.1
    TREND_WINDOW = 20

    daily_return = (self.portfolio_value -

```

```

        self.previous_portfolio) / self.previous_portfolio

self.returns.append(daily_return)
self.max_portfolio = max(self.max_portfolio, self.portfolio_value)

reward = 0

curr_vix = self._get_vix()

drawdown_threshold = curr_vix * 0.005

drawdown = (self.max_portfolio - self.portfolio_value) / \
    self.max_portfolio

if drawdown > drawdown_threshold:
    reward -= drawdown

if len(self.returns) >= SHARPE_WINDOW_SHORT:
    ret = np.array(self.returns[-SHARPE_WINDOW_SHORT:])
    sharpe_ratio = np.sqrt(SHARPE_WINDOW_SHORT) * np.mean(ret) / np.std(ret)

    reward += sharpe_ratio

if len(self.returns) >= SHARPE_WINDOW_MED:
    ret = np.array(self.returns[-SHARPE_WINDOW_MED:])
    sharpe_ratio = np.sqrt(SHARPE_WINDOW_MED) * np.mean(ret) / np.std(ret)

    reward += sharpe_ratio * 0.5

if len(self.returns) >= SHARPE_WINDOW_LONG:
    ret = np.array(self.returns[-SHARPE_WINDOW_LONG:])
    sharpe_ratio = np.sqrt(SHARPE_WINDOW_LONG) * np.mean(ret) / np.std(ret)

    reward += sharpe_ratio * 0.25

if len(self.returns) > 0:
    curr_volatility = np.std(
        np.array(self.returns[-min(SHARPE_WINDOW_SHORT, len(self.returns)):]))

    volatility_weight = curr_vix * 0.01

    volatility_penalty = curr_volatility * volatility_weight

    reward -= volatility_penalty

if self.curr_step > 0:
    curr_date = self.trading_days[self.curr_step - 1]

    for ii, action in enumerate(action_list):

        if self.curr_step >= SHARPE_WINDOW_SHORT:
            prev_date = self.trading_days[self.curr_step - SHARPE_WINDOW_SHORT]
            stock_price = self.stock_data.xs(
                self.ticker_list[ii], level="Ticker", axis=0).loc[prev_date:curr_date, "Adj Close"].values

            if len(stock_price) > 0:
                ema = np.mean(stock_price)

                curr_price = self.stock_data.loc[(
                    self.ticker_list[ii], curr_date), "Adj Close"]

                if len(stock_price) >= TREND_WINDOW:
                    returns = np.diff(stock_price) / stock_price[:-1]
                    volatility = np.std(returns[-TREND_WINDOW:])
                    trend_reward = BASE_TREND_REWARD * (1 + volatility)
                else:
                    trend_reward = BASE_TREND_REWARD

                if (action > 0 and curr_price > ema) or (action < 0 and curr_price < ema):
                    reward += trend_reward * abs(action)

    return reward

def _get_observation_ryan_0(self):
    """
    Returns observations for current time step:

    Parameters:

    Returns:
        observation (list): The owned shares, current and previous opening
        prices, and volume for each stock in the portfolio
    """

    curr_date = self.trading_days[self.curr_step - 1]
    if self.curr_step != 1:
        prev_date = self.trading_days[self.curr_step - 2]

```

```

else:
    prev_date = self.trading_days[self.curr_step - 1]

opening_price = np.array([self.stock_data.loc[(
    ticker, curr_date), "Open"] for ticker in self.ticker_list])
opening_price_prev = np.array([self.stock_data.loc[(
    ticker, prev_date), "Open"] for ticker in
    self.ticker_list])

volume = np.array([self.stock_data.loc[(ticker, curr_date), "Volume"]
    for ticker in self.ticker_list])

observation = [round(self.curr_funds, 3)]
for ii in range(self.num_stocks):
    observation.append(int(self.owned_shares[ii]))
    observation.append(round(opening_price[ii], 3))
    observation.append(round(opening_price_prev[ii], 3))
    observation.append(round(volume[ii], 3))

return observation

def _get_vix(self):
    """
    Retrieves the VIX index for the current trading day:

    Parameters:

    Returns:
        curr_vix (float): The VIX opening price for the current trading day
    """

    curr_vix = self.vix["Open"].iloc[-1]
    return curr_vix

```